True or false

Part 1

1. An atlas 900 : True
2. Ascend ai : True
3. In actual data processing : true
4. The modelArts : true
5. Filter methods : false
6. A decision tree : false
7. Object storage service : false
8. In the convolutional layer, : true
9. During model training : false
10. Tensorflow 2.x : false
11. An ai processor : true
12. Creative work, such : true
13. modelArts Allows : false
14. Mindspore accelerates : True
15. Convolutional neural networks ; false
16. The xor problems : true
17. Mini-batch : false
18. As the cornerstone of : false
19. In the deep neural networks : false
20. Poisson distribution is : false

Part 2

1. In actual data processing : true
2. Mindspore supports automatic : false
3. Dropout randomly discards : false
4. Boosting ; false
5. Although each : true
6. The cost of L1 : true
7. Assuming the number : false
8. If a 32x32 : false
9. Tensorflow is Microsoft : false
10. Mindspore accelerates : true
11. Ascend Ai : true
12. The modelarts : true
13. The dropout technology : false
14. Powered by ascend AI processors : true
15. The PIL module in tensorflow 2. : false
16. Assuming dataset contains : false
17. Mini-batch gradients : false
18. An atlas 900 AI : true
19. Poisson distribution : false
20. Creative work : true

Part 3 :

1. as the cornerstone of the Huawei : false
2. all kernels of the same : true
3. an ai processor is also : true
4. device, edge ,cloud, : true
5. convolutional neural : false
6. creative work : true
7. object storage service : false
8. Huawei hiAI 3.0
9. The tree model used : false
10. Tensorflow 2.x is : false
11. Tensorflow is microsft : false
12. The development of ai : true
13. Federated learning : true
14. In industry application : false
15. When collecting images : true
16. All neurons : true
17. Dropout randomly : false
18. Ascend AI : true
19. Assuming the number : false

PART 4 :

A model constructed using : true

Device,edge : true

With the continued development : true

The softmax function turns : false

In actual data : true

All kernels of the same : false

Mindspore supports : true

In the deep neural networks : false

The spatial relationship : true

As the cornerstone of : false

The PIL module : false

An ai processor : true

When collecting image : false

Mini-batch gradient : false

The added cost of : true

Part 5

1. In the deep neural : false
2. The tensor boost engine (TBE): false
3. An ai processor is also : true
4. In the convolutional layer : true
5. Powered by ascend AI : false
6. All kernels of the same : False’
7. Huawei HiAI 3.0 : true
8. Mindspore accelerates model : true
9. Convolutional neural : false
10. In tensorflow 2.0
11. Although each industry and : true
12. Typically, a larger : true
13. In industry application, models : false
14. Device, edge : true
15. The development of ai : true
16. Tensorflow is a Microsoft :false
17. Objext storage service (OBS) : false
18. An atlas 900 : true
19. Ascend AI processor : true
20. The cost L1 : true

SINGLE ANSWER

Part 1

21. Which of the following is the main function of HUAWEI CLOUD GeoGenius?

A. Ecological environment detection

C. Real-time traffic detection

D. Traffic incident detection

B. Heat demand forecast

Congratulations! Correct Answer: A

22. Which of the following is not a feature of the MindSpore core architecture?

B. Automatic parallelism

C. Automatic deployment

A. Automatic differentiation

D. Automatic tuning

Congratulations! Correct Answer: C

23. Which one of the following actions can the function tf.squeeze be used for in TensorFlow 2.0?

C. Tensor concatenation

B. Computation of the absolute value

A. Element-wise addition

D. Dimensionality reduction

Congratulations! Correct Answer: D

24. Which of the following statements is false about the ReLU function?

B. The ReLU function effectively alleviates the vanishing gradient problem.

A. The ReLU function is not differentiable at x = 0 and a derivative is forcibly defined at this point.

D. Compared with Sigmoid and tanh, the convergence of the ReLU function is slow.

C. The surface defined at the zero point of the ReLU function is not smooth enough in some regression problems.

Congratulations! Correct Answer: D

25. Which of the following is a lightweight and high-performance service module that helps MindSpore developers efficiently deploy online inference services in production environments?

A. MindIR

C. MindArmour

B. MindSpore Serving

D. MindInsight

Congratulations! Correct Answer: B

26. Which of the following functions is used in HUAWEI CLOUD General Text OCR experiments?

B. ocr_client.ocr_service_base64

D. ocr.request_ocr_service_base128

A. ocr_client.request_ocr_service_base64

C. ocr_client.request_base64

Congratulations! Correct Answer: A

27. Which of the following indicators cannot be used to evaluate a model?

D. Code complexity

C. Prediction rate

B. Explainability

A. Generalization capability

Congratulations! Correct Answer: D

28. Dirty data refers to data with quality problems. Which of the following statements is false about the data quality?

D. Inconsistency: Data contains conflicting records.

A. Noise: Data contains incorrect records or exceptions.

C. Unprocessed: Data for which feature engineering has not been performed.

B. Incompleteness: Data lacks attributes or contains some missing values.

Congratulations! Correct Answer: C

29. concat sample 1 tf random normal ([4,100,100,3])</br>concat_sample 2-tf.random.normal([40,100,100,3])</br>concated sample_1 = tf.concat([concat_sample_1.concat_sample_2],axis=0)</br> print(concated sample_1.shape) </br>Which of the following is the output for this code?

A. (44,200,200,6)

B. (44,100,100,3)

C. (40,100,100,6)

D. (44,100,100,6)

Congratulations! Correct Answer: B

30. Which of the following is NOT an application of speech processing?

C. Voiceprint recognition

A. Speech recognition

B. Text to speech

D. Named entity recognition

Congratulations! Correct Answer: D

31. Which of the following is the type of labels predicted by ensemble learning algorithms?

D. Strain and discrete

A. Discrete and continuous

C. Discrete

B. Continuous

Congratulations! Correct Answer: A

32. Which operation is not a step in the network definition process during application development?

C. Defining an optimizer and a loss function

A. Specifying an input dimension

D. Model compression

B. Weight initialization

Sorry! Correct Answer: D

33. Which of the following determines the upper limit of machine learning?

C. Scenarios

A. Data

D. Computing power

B. Algorithms

Congratulations! Correct Answer: A

34. Intelligent quality inspection is based on the cloud-edge-device synergy of deep learning algorithms. Which of the following operations is performed on edge devices?

C. Model delivery

A. Model training

B. Data labeling

D. Onsite inference

Congratulations! Correct Answer: D

35. A vendor wants to provide an intelligent EMR system for a hospital Which of the following technologies is involved in the system?

C. Expert systems and knowledge graph

A. Natural language processing

D. Image generation and enhancement

B. Object detection

Sorry! Correct Answer: A

36. Which of the following statements is false about gradient descent algorithms?

B. When there are too many samples and GPUs are not used for parallel computing, the convergence process of the global gradient descent is time-consuming.

A. The global gradient descent is more stable than the stochastic gradient descent (SGD).

C. When GPUs are not used for parallel computing, the mini-batch gradient descent (MBGD) takes less time than the SGD to complete an epoch.

D. Fach time the global descent updates its weight, all training samples need to be calculated.

Congratulations! Correct Answer: C

37. Which of the following statements is false about machine learning?

A. Model generalization capability: the extent to which a learned model can be applied to new samples, also called robustness.

B. Error: difference between the prediction of a learned model on a sample and the actual result of the sample. Errors can be classified into training errors and generalization errors.

C. Underfitting: The training error is large.

D. Overfitting: The training error is large.

Congratulations! Correct Answer: D

38. In the Da Vinci Architecture, which of the following computation data types is supported by the vector unit?

C. FP32

B. BFloat16

D. BFloat32

A. INT16

Congratulations! Correct Answer: C

Part 2

21. Which of the following statements is true about optimizers?

D. The learning rate of the momentum optimizer does not need to be manually set.

B. Unlike the RMS prop optimizer, the Adagrad optimizer does not automatically update the learning rate.

C. The SGD and momentum optimizers use the same learning rate for each iteration.

A. The RMSprop optimizer inherits the advantages of the Adam optimizer. The learning rate is automatically updated.

Congratulations! Correct Answer: C

22. The Relu function is often used in deep learning and neural networks. Which of the following statements is true about this function?

A. The value range is [-1,1].

D. The value range of the derivatives is [-1,0].

B. The value range is [0.+\).

C. The value range of the derivatives is (0,1).

Congratulations! Correct Answer: B

23. Which of the following statements is false about decision trees?

D. The variance can be regarded as a quantitative index of purity.

C. The key step to building a decision tree involves splitting it based on feature attributes.

B. The C4.5 algorithm uses the information gain ratio to select feature attributes.

A. Except the root node, all nodes in a decision tree are called leaf nodes.

Congratulations! Correct Answer: A

24. Which of the following functions is used in HUAWEI CLOUD General Text OCR experiments?

B. ocr_client.ocr_service_base64

D. ocr.request_ocr_service_base128

A. ocr_client.request_ocr_service_base64

C. ocr_client.request_base64

Congratulations! Correct Answer: A

25. Which of the following statements is true about model validity?

B. Variance refers to the difference between the prediction of a learned model on a sample and the actual result of the sample.

D. Bias refers to errors you get when you run the model on new samples.

C. During model testing, errors can be classified into training errors and sample variances.

A. The goal of machine learning is for a trained model to perform well on new samples, not just on samples used for training.

Congratulations! Correct Answer: A

26. Which of the following statements is true about variance and bias?

C. A model with low bias and variance has low precision.

D. A model with high bias and variance performs poorly and will not be used.

A. A model with low bias but high variance has high robustness.

B. A model with high bias but low variance has high precision.

Congratulations! Correct Answer: D

27. Which of the following is the correct shape of tensor [[0,1], [2,3], [4,5]]?

C. (1,6)

A. (2,3)

D. (6,1)

B. (3,2)

Congratulations! Correct Answer: B

28, HUAWEI CLOUD Traffic Intelligent Twins (or TrafficGo) is a comprehensive urban traffic management solution. Which of the following functions is provided by this solution?

D. Assisted autonomous driving

B. Automatic path planning

C. Emergency response and disaster prevention

A. Traffic parameter awareness

Congratulations! Correct Answer: A

29. Which of the following statements is true about neural networks?

A. The full name of LSTM is Less Short-Term Memory.

B. LSTM has nothing to do with RNNs.

D. Neural networks can only propagate information forward.

C. LSTM is suitable for processing events with long interval and delay in the time sequence.

Congratulations! Correct Answer: C

30. Which of the following is not a feature of HIAI 2.0?

D. Distributed

A. Single-device

C. Device-cloud synergy

B. Multi-device

Congratulations! Correct Answer: D

31. A vendor wants to provide an intelligent EMR system for a hospital. Which of the following technologies is involved in the system?

A. Natural language processing

C. Expert systems and knowledge graph

B. Object detection

D. Image generation and enhancement

Congratulations! Correct Answer: A

32. Which of the following is not supported by data management on ModelArts?

B. No-manual labeling

A. A wide range of data formats

D. Intelligent data analysis

C. Intelligent data filtering

Congratulations! Correct Answer: B

33. Which of the following is covered by HUAWEI CLOUD EI Health?

C. Patient care

D. Doctor-patient dispute resolution

A. Genome

B. Smart surgery

Congratulations! Correct Answer: A

34. A generative adversarial network is like a game system. The generator produces fake samples, and the discriminator tries to distinguish real data from the data created by the generator. What is the ideal result?

C. The discriminator distinguishes real data from the data created by the generator.

A. The generator is trained to fool the discriminator.

B. The discriminator cannot distinguish real data from the data created by the generator.

D. Develop a high-precision discriminator and a generator that cannot fool the discriminator.

Congratulations! Correct Answer: B

35. Which of the following is NOT a main function of HUAWEI CLOUD GeoGenius?

A. Natural resource survey

B. Ecological environment detection

D. Real-time traffic detection

C. Weather forecast

Sorry! Correct Answer: D

36. When the v1 compatibility package of TensorFlow 2.0 is used to inherit the Tensorflow 1.x code, the eager operation needs to be disabled. Which of the following commands is used to disable it?

C. tf.uneager_execution()

B. tf.no_eager_execution()

A. tf.disable_eager_run()

D. tf.disable_eager_execution()

Sorry I Correct Answer: D

37. In the Da Vinci Architecture, which of the following computation data types is supported by the vector unit?

B. BFloat16

A. INT16

D. BFloat32

C. FP32

Congratulations! Correct Answer: C

38. Which of the following statement is false about the running process of the MindArmour subsystem?

D. Trustworthiness enhancement: Use preset methods to enhance the trustworthiness of Al models.

B. Fuzzing execution: Generate trusted test data randomly based on the model coverage and configuration policies.

C. Evaluation report: Generate an evaluation report based on built-in or user-defined trustworthiness metrics.

A. Configuration policies: Define test policies based on threat vectors and trustworthiness certification requirements and select appropriate test data generation methods.

Congratulations! Correct Answer: B

Part 3

21. Which of the following is the reasoning method of the production system that draws conclusions through a rule library?

B. Backward

D. Random

C. Bidirectional

A. Forward

Congratulations! Correct Answer: A

22. Which of the following is a math operator in MindSpore?

D. Select

A. Mul

C. Callback

B. ControlDepend

Congratulations! Correct Answer: A

23. Which of the following is not a hyperparameter for model training?

C. b in the one-dimensional linear regression model Y=wX+b

D. Number of trees in a random forest

B. K in K-means

A. batch_size in a training neural network

Congratulations! Correct Answer: C

24. HUAWEI CLOUD Traffic Intelligent Twins (or TrafficGo) is a comprehensive urban traffic management solution. Which of the following functions is provided by this solution?

B. Automatic path planning

D. Assisted autonomous driving

A. Traffic parameter awareness

C. Emergency response and disaster prevention

Congratulations! Correct Answer: A

25. Which operation is not a step in the network definition process during application development?

A. Specifying an input dimension

B. Weight initialization

D. Model compression

C. Defining an optimizer and a loss function

Sorry! Correct Answer: D

26. Which of the following products can be used by a company that wants to provide intelligent customer

service?

A. RDB

D. GES

C. OBS

B. Conversational Bot

Congratulations! Correct Answer: B

27. A generative adversarial network is like a game system. The generator produces fake samples, and the discriminator tries to distinguish real data from the data created by the generator. What is the ideal result?

D. Develop a high-precision discriminator and a generator that cannot fool the discriminator.

A. The generator is trained to fool the discriminator.

B. The discriminator cannot distinguish real data from the data created by the generator.

C. The discriminator distinguishes real data from the data created by the generator.

Congratulations! Correct Answer: B

28. Which of the following is NOT an application of speech processing?

B. Text to speech

C. Voiceprint recognition

D. Named entity recognition

A. Speech recognition

Congratulations! Correct Answer: D

29. Which of the following is not a feature of HiAI 2.0?

A. Single-device

D. Distributed

B. Multi-device

C. Device-cloud synergy

Congratulations! Correct Answer: D

30. Which of the following statements is true about the loss functions typically used in deep learning?

B. The quadratic cost function reflects the gap between the target output and the actual output.

A. Quadratic cost functions are usually used for classification, while cross-entropy cost functions are used for clustering.

C. The quadratic cost function reflects the gap between two probability distributions.

D. The purpose of training is to minimize the loss function.

Congratulations! Correct Answer: B

31. Which of the following is a method of compressing an Al model?

C. Data augmentation

A. Low rank approximation

D. Data volume reduction

B. Network integration

Congratulations! Correct Answer: A

32. Which of the following statements is true about variance and bias?

A. A model with low bias but high variance has high robustness.

C. A model with low bias and variance has low precision.

D. A model with high bias and variance performs poorly and will not be used.

B. A model with high bias but low variance has high precision.

Congratulations! Correct Answer: D

33. Which of the following is the purpose of running ocr_client = HWOcrClientAKSK(ak, sk, region) when using HUAWEI CLOUD OCR?

B. Transfer images for recognition.

D. Encode image data.

A. Initialize the client.

C. Delete image data.

Congratulations! Correct Answer: A

34. How does Al perform in the cognitive intelligence phase?

C. Machines begin to understand, think, and make decisions like humans.

A. Machines begin to compute and transmit information like humans.

D. Machines begin to record information, such as pictures and voices.

B. Machines begin to see, listen, understand, make judgments, and take simple actions.

Sorry! Correct Answer: C

35. Which of the following is the correct shape of tensor [[[0,1],[2,3]], [[4,5],[6,7]]]?

A. [3,3,2]

C. [2,3,4]

B. [3,2,4]

D. [2.2.2]

Congratulations! Correct Answer: D

36. Which of the following statements is true about the ReLU function?

C. There is no upper bound, so training is relatively easy to diverge.

B. Neuron necrosis does not occur.

A. The ReLU function is differentiable at x = 0, and the derivative is 1.

D. The gradient is always 1, and the vanishing gradient problem can be perfectly solved.

Congratulations! Correct Answer: C

37. Which of the following is NOT a main function of HUAWEI CLOUD GeoGenius?

C. Weather forecast

D. Real-time traffic detection

B. Ecological environment detection

A. Natural resource survey

Congratulations! Correct Answer: D

38. When the v1 compatibility package of TensorFlow 2.0 is used to inherit the TensorFlow 1.x code, the eager operation needs to be disabled. Which of the following commands is used to disable it?

A. tf.disable_eager_run()

B. tf.no_eager_execution()

C. tf.uneager_execution()

D. tf.disable_eager_execution()

Congratulations! Correct Answer: D

Part 4

21. Which one of the following actions can the function tf.squeeze be used for in TensorFlow 2.0?

C. Tensor concatenation

D. Dimensionality reduction

A. Element-wise addition

B. Computation of the absolute value

Congratulations! Correct Answer: D

22, Which of the following products can be used by a company that wants to provide intelligent customer service?

D. GES

C. OBS

A. RDB

B. Conversational Bot

Congratulations! Correct Answer: B

23. How does Al perform in the cognitive intelligence phase?

A. Machines begin to compute and transmit information like humans.

C. Machines begin to understand, think, and make decisions like humans.

B. Machines begin to see, listen, understand, make judgments, and take simple actions.

D. Machines begin to record information, such as pictures and voices.

Congratulations! Correct Answer: C

24, Which operation is not a step in the network definition process during application development?

A. Specifying an input dimension

C. Defining an optimizer and a loss function

D. Model compression

B. Weight initialization

Congratulations! Correct Answer: D

25. Which of the following commands is used to install TensorFlow?

C. conda setup.py tensorflow

D. pip install tensorflow

B. python tensorflow setup.py

A. conda create install tensorflow

Congratulations! Correct Answer: D

26. Which of the following is the type of labels predicted by ensemble learning algorithms?

D. Strain and discrete

C. Discrete

A. Discrete and continuous

B. Continuous

Congratulations! Correct Answer: A

27, Which of the following HUAWEI CLOUD services is NOT a basic Al platform service?

D. Graph Engine Service (GES)

C. Relational Database Service (RDS)

B. Huawei HiLens

A. ModelArts

Congratulations! Correct Answer: C

28. An image recognition experiment uses 42,510 training images. There are more than 200 categories, each of which contains only 10 images. There are also categories with only 20 to 50 images. Which of the following data problems does this phenomenon fit into?

B. Data imbalance

C. Data loss

A. Data augmentation

D. Data fitting

Congratulations! Correct Answer: B

29. Which of the following statements is false about the ReLU function?

C. The surface defined at the zero point of the ReLU function is not smooth enough in some regression problems.

D. Compared with Sigmoid and tanh, the convergence of the ReLU function is slow.

A. The ReLU function is not differentiable at x = 0 and a derivative is forcibly defined at this point.

B. The ReLU function effectively alleviates the vanishing gradient problem.

Congratulations! Correct Answer: D

30. Which of the following is not a common tensor operation in MindSpore?

D. isinstance0

A. asnumpy0

B. size0

C. dim0

Congratulations! Correct Answer: D

31. In the Da Vinci Architecture, which of the following computation data types is supported by the vector unit?

A. INT16

D. BFloat32

C. FP32

B. BFloat16

Congratulations! Correct Answer: C

32. Which of the following statement is false about the running process of the MindArmour subsystem?

A. Configuration policies: Define test policies based on threat vectors and trustworthiness certification requirements and select appropriate test data generation methods.

B. Fuzzing execution: Generate trusted test data randomly based on the model coverage and configuration policies.

D. Trustworthiness enhancement: Use preset methods to enhance the trustworthiness of Al models.

C. Evaluation report: Generate an evaluation report based on built-in or user-defined trustworthiness metrics.

Congratulations!Correct Answer: B

33. Which of the following indicators cannot be used to evaluate a model?

D. Code complexity

C. Prediction rate

A. Generalization capability

B. Explainability

Congratulations! Correct Answer: D

34. Which of the following statements is false about machine learning?

A. Model generalization capability: the extent to which a learned model can be applied to new samples, also called robustness.

C. Underfitting: The training error is large.

B. Error: difference between the prediction of a learned model on a sample and the actual result of the sample. Errors can be classified into training errors and generalization errors.

D. Overfitting: The training error is large.

Congratulations!Correct Answer: D

35. Which of the following statements is true about deep learning?

C. Parameters at each layer have specific meanings.

B. GPU parallel computing provides faster computing if there are more network parameters.

A. Algorithms are used to automatically extract features.

D. The feature interpretability is strong.

Congratulations!Correct Answer: A

36. The Relu function is often used in deep learning and neural networks. Which of the following statements is true about this function?

B. The value range is [0,+\).

C. The value range of the derivatives is (0,1).

A. The value range is [-1,1].

D. The value range of the derivatives is [-1,0].

Congratulations! Correct Answer: B

37. In TensorFlow 2.0, x = tf.constant([1,2,3])</br>y = tf.broadcast_to(x, [3, 3])</br>
print(y)</br>Which of the following is the output for this code?

B. [[1,2,3], [1,2,3], [1,2,3]]

D. [[1,1,1], [2,2,2), (3.3.3]]

A. [[1,2,3]]

C. [[1,2,3,1,2,3,1,2,3]]

Congratulations! Correct Answer: B

38. Which of the following statements is true about variance and bias?

D. A model with high bias and variance performs poorly and will not be used.

A. A model with low bias but high variance has high robustness.

C. A model with low bias and variance has low precision.

B. A model with high bias but low variance has high precision.

Congratulations! Correct Answer: D

Part 5

21. Which of the following is a lightweight and high-performance service module that helps MindSpore developers efficiently deploy online inference services in production environments?

D. MindInsight

C. MindArmour

B. MindSpore Serving

A. MindIR

22. Which of the following statements is false about machine learning?

C. Underfitting: The training error is large.

B. Error: difference between the prediction of a learned model on a sample and the actual result of the sample. Errors can be classified into training errors and generalization errors.

D. Overfitting: The training error is large.

A. Model generalization capability: the extent to which a learned model can be applied to new samples, also called robustness.

23. Which of the following statements is false about gradient descent algorithms?

D. Each time the global descent updates its weight, all training samples need to be calculated.

C. When GPUs are not used for parallel computing, the mini-batch gradient descent (MBGD) takes less time than the SGD to complete an epoch.

B. When there are too many samples and GPUs are not used for parallel computing, the convergence process of the global gradient descent is time-consuming.

A. The global gradient descent is more stable than the stochastic gradient descent (SGD).

24. An image recognition experiment uses 42,510 training images. There are more than 200 categories, each of which contains only 10 images. There are also categories with only 20 to 50 images. Which of the following data problems does this phenomenon fit into?

B. Data imbalance

C. Data loss

D. Data fitting

A. Data augmentation

25. "Knowledge representation is the unique method of representing knowledge using a set of symbols in a structure that can be understood by computers." Which of the following is true about this statement?

A. This statement is false. The method of representing knowledge is not unique.

C. This statement is correct. The knowledge representation can support expert systems.

D. This statement is false. Knowledge representation cannot be used for expert rules or fuzzy inference.

B. This statement is correct. The coding method of knowledge representation is unique.

26. Which of the following is the correct shape of tensor [[[0, 1],[2,3]], [[4,5],[6,7]]]?

A. [3,3,2]

B. [3,2,4]

C. [2,3,4]

D. [2,2,2]

27. Which of the following statements is true about deep learning?

C. Parameters at each layer have specific meanings.

D. The feature interpretability is strong.

B. GPU parallel computing provides faster computing if there are more network parameters.

A. Algorithms are used to automatically extract features.

28. Which of the following statements is true about rule-based approaches?

D. Decision-making rules are complex or difficult to describe.

B. Rule-based approaches use explicit programming to solve problems.

C. Rule-based approaches use samples for training.

A. Rules are obtained through machine learning.

29. Which of the following is not a feature of the MindSpore core architecture?

A. Automatic differentiation

B. Automatic parallelism

C. Automatic deployment

D. Automatic tuning

30. Which of the following is a math operator in MindSpore?

A. Mul

B. ControlDepend

C. Callback

D. Select

31. Which of the following statements is true about the loss functions typically used in deep learning?

C. The quadratic cost function reflects the gap between two probability distributions.

A. Quadratic cost functions are usually used for classification, while cross-entropy cost functions are used for clustering.

D. The purpose of training is to minimize the loss function.

B. The quadratic cost function reflects the gap between the target output and the actual output

32. Which of the following statements is true about model validity?

C. During model testing, errors can be classified into training errors and sample variances.

D. Bias refers to errors you get when you run the model on new samples.

B. Variance refers to the difference between the prediction of a learned model on a sample and the actual result of the sample.

A. The goal of machine learning is for a trained model to perform well on new samples, not just on samples used for training.

33. "Deep learning is a complex, comprehensive discipline. It includes Al and machine learning." Which of the following is true about this statement?

A. This statement is correct. Deep learning includes Al.

D. This statement is incorrect. Al includes deep learning and machine learning.

C. This statement is incorrect. Machine learning has nothing to do with deep learning.

B. This statement is correct. Deep learning includes machine learning.

34. Which of the following statements is false about decision trees?

B. The C4.5 algorithm uses the information gain ratio to select feature attributes.

A. Except the root node, all nodes in a decision tree are called leaf nodes.

D. The variance can be regarded as a quantitative index of purity.

C. The key step to building a decision tree involves splitting it based on feature attributes.

35. Which of the following is the main function of HUAWEI CLOUD GeoGenius?

B. Heat demand forecast

C. Real-time traffic detection

D. Traffic incident detection

A. Ecological environment detection

36. Which of the following statement is false about the running process of the MindArmour subsystem?

B. Fuzzing execution: Generate trusted test data randomly based on the model coverage and configuratio policies.

C. Evaluation report: Generate an evaluation report based on built-in or user-defined trustworthiness metrics.

D. Trustworthiness enhancement: Use preset methods to enhance the trustworthiness of Al models.

A. Configuration policies: Define test policies based on threat vectors and trustworthiness certification requirements and select appropriate test data generation methods.

37. Which of the following commands is used to install TensorFlow?

D. pip install tensorflow

C. conda setup.py tensorflow

B. python tensorflow setup.py

A. conda create install tensorflow

38. The core of an Al framework and one of the decisive factors of the programming paradigm is the automatic differentiation used by the Al framework. Which of the following is used by

MindSpore?

B. Operator overloading

D. Just-in-time compilation

C. Source to Source (525)

A. Graph kernel fusion

MULTI ANSWER

PART 1

39. Which of the following are common tensor operations in MindSpore?

D. abs0

B. astype()

C. switch()

A. dtype()

Congratulations! Correct Answer: ABD

40. Which of the following are part of the Huawei full-stack AI solution?

B. openEuler

A. TBE

C. AscendCL

D. Ascend

Congratulations! Correct Answer: ACD

41. Which of the following are the features of the Speech Interaction Service (SIS) provided by HUAWEI CLOUD EI?

A. Multiple (more than three) languages

C. Low requirements

B. High usability

D. Low labor cost

Congratulations! Correct Answer: BCD

42. Which of the following statements are true about MindSpore and Huawei all-scenario solutions?

C. Distributed training of ultra-large models and ultra-large datasets requires only data parallelism.

D. MindIR is the key to "Train Once, Deploy Everywhere."

B. MindSpore Serving supports batching.

A. The full-scenario deployment solution includes model generation and efficient execution.

Congratulations! Correct Answer: ABD

43. Which of the following are CNN-based network structures?

C. Inception

B. Transformer

D. BERT

A. ResNet

Congratulations! Correct Answer: AC

44. What are the theories of cognitive intelligence?

A. Threshold theory

C. Behavioral theory

D. Logic theory

B. The theory of evolution

Congratulations! Correct Answer: ABD

45. Which of the following are used for backpropagation?

B. Activation functions

A. Number of network layers

C. Number of neurons at a single layer

D. Neuron weight

Congratulations! Correct Answer: BD

46. Which of the following are major advantages of Keras?

B. Low-level coding style

C. Modular and combinable

D. Easy to scale

A. Easy to use

Congratulations! Correct Answer: ACD

47. Which of the following are model hyperparameters?

D. Weight coefficient w in linear regression

A. Learning rate, iteration count, and batch size in a training neural network

B. K in the k-nearest neighbors (KNN) algorithm

C. Number of trees in a random forest

Congratulations! Correct Answer: ABC

48. Which of the following statements are true about linear regression?

D. Due to algorithm complexity, linear regression cannot use the gradient descent method to calculate the weight parameter if the loss function reaches the minimum value.

B. The loss function of linear regression can be obtained using the normal distribution function and maximum likelihood estimation (MLE).

C. The result of multivariate linear regression analysis may be a hyperplane in a high-dimensional space.

A. The error of linear regression is affected by many independent factors. According to the central limit theorem (CLT), the error follows normal distribution.

Congratulations! Correct Answer: ABC

49. Which of the following statements are true about ensemble learning?

C. A batch of features are randomly selected for the subtree training in a random forest.

D. Xgboost supports parallel training like the random forest.

A. Ensemble learning uses the decision tree as the bottom-layer algorithm.

B. The voting policy of ensemble learning can be the averaging method.

Congratulations! Correct Answer: BC

50. Which of the following statements are false about Keras?

B. Keras is a neural network development package used to build CNN sequential models. It cannot be used to build other neural networks.

A. Like TensorFlow, Keras is a multi-layer neural network development package. However, Keras has simpler syntax and is easier to use.

C. The neural network model built using Keras must be compiled before data can be input into it for training.

D. Keras can be used as the backend of TensorFlow.

Congratulations! Correct Answer: BD

51. Which of the following functions are supported by the HIAI Engine platform?

A. Form recognition

C. Video summarization

B. Keyword extraction

D. Automatic generation of video thumbnails

Congratulations! Correct Answer: ABCD

52. Which of the following cells are predefined in MindSpore?

D. WithGradCell

C. cos

OB. TFRecord

A. MAELoss

Congratulations! Correct Answer: AD

53. Which of the following are common ensemble learning algorithms in machine learning?

B. GBDT

A. Random forest

C. Polynomial regression

D. Adaboost

Congratulations! Correct Answer: ABD

54. Which of the following statements are false about common optimizers?

C. Adam sets different learning rates for different parameters.

D. The RMSprop optimizer introduces an attenuation coefficient to ensure a certain proportion of gradient attenuation in each round.

B. When the momentum optimizer is used, parameters are updated by using the same learning rate, but momentum coefficients keep changing with each iteration.

A. One of the advantages of Adagrad optimizers is that the parameter update operation does not end too early.

Congratulations! Correct Answer: AB

PART 2

39. What are the theories of cognitive intelligence?

C. Behavioral theory

D. Logic theory

B. The theory of evolution

A. Threshold theory

Congratulations! Correct Answer: ABD

40. Which of the following statements are true about decision trees?

D. Building a decision tree means selecting and measuring feature attributes and determining their topology.

A. A decision tree is a binary tree structure.

B. Each non-leaf node in a decision tree represents a test on a feature attribute; each branch represents the output of the feature attribute within a value range; and each leaf node holds a class label.

C. Except the root node, all nodes in a decision tree are called leaf nodes.

Congratulations! Correct Answer: BD

41. Which of the following are CNN-based network structures?

C. Inception

B. Transformer

D. BERT

A. ResNet

Congratulations! Correct Answer: AC

42. Which of the following technologies may be involved in room service robots?

B. Speech recognition

D. Object detection

A. Path planning

C. Sentiment analysis

Congratulations! Correct Answer: ABD

43. Which of the following methods provided by TensorFlow 2.0 cannot be used to check whether an object is a tensor?

C. device

A. is tensor

D. tftypes

B. isinstance

Congratulations! Correct Answer: BCD

44. Which of the following functions are supported by the HiAl Engine platform?

A. Form recognition

D. Automatic generation of video thumbnails

B. Keyword extraction

C. Video summarization

Congratulations! Correct Answer: ABCD

45. Which of the following operations are involved when TensorFlow is used for model training?

B. Train and save models.

D. Define learning parameters and operations.

C. Obtain and cleanse data.

A. Define input and output nodes.

Congratulations! Correct Answer: ABD

46. Which of the following statements are true about grid search based on hyperparameter tuning?

C. Grid search works well when there are relatively few hyperparameters.

D. Grid search suits neural networks well.

B. Grid search is expensive and time-consuming.

A. Grid search exhaustively searches for all possible hyperparameter combinations to form a hyperparameter value grid.

Congratulations! Correct Answer: ABC

47. Which of the following use the Da Vinci Architecture?

B. RTX3080

C. Kunpeng 920

D. Ascend 910

A. Ascend 310

Congratulations! Correct Answer: AD

48. Which of the following statements are false about the functions of the pooling layers?

B. Max pooling achieves invariance to some extent.

D. Pooling helps prevent overfitting.

C. Pooling reduces the size of the input data at the next layer and the number of parameters, but increases the computation amount.

A. They reduce the receptive field.

Congratulations! Correct Answer: AC

49. Which of the following statements are true about linear regression?

D. Due to algorithm complexity, linear regression cannot use the gradient descent method to calculate the weight parameter if the loss function reaches the minimum value.

C. The result of multivariate linear regression analysis may be a hyperplane in a high-dimensional space.

A. The error of linear regression is affected by many independent factors. According to the central limit theorem (CLT), the error follows normal distribution.

B. The loss function of linear regression can be obtained using the normal distribution function and maximum likelihood estimation (MLE).

Congratulations! Correct Answer: ABC

50. Which of the following are common application scenarios of HUAWEI CLOUD OCR?

C. VAT invoice recognition

D. Aircraft itinerary recognition

A. Photo recognition for cats and dogs

B. Driving license recognition

Congratulations! Correct Answer: BCD

51. Which of the following statements are true about activation functions?

C. There are many activation functions. You need to select one based on the actual situation.

B. If we do not use an activation function, the output signals will be just a simple linear function.

A. The existence of the activation function introduces linearity into the network.

D. Activation functions play an important role in learning and understanding complex and nonlinear functions of neural network models.

Congratulations! Correct Answer: BCD

52. Which of the following statements are true about the MindSpore components?

C. nn: neural network cells in MindSpore that define loss functions and optimizers.

A. communication: processes data flow communication between the CPU and memory.

D. train: relates to training model and model quantization module.

B. dataset: dataset processing module that loads and preprocesses data.

Sorry! Correct Answer: BC

53. Backed by HUAWEI CLOUD's accumulated knowledge and expertise in Al, big data, and other cutting-edge technologies, GeoGenius offers a one-stop Al development cloud platform for remote sensing data processing, mining, and management. Which of the following are the main functions of GeoGenius?

A. Weather forecasting

C. Emergency response and disaster prevention

B. Afforestation

D. Smart transport infrastructure

Sorry! Correct Answer: AC

54. Which of the following statements are false about GANS?

A. The discriminator input is mainly noise and real sample data.

B. GANS are a type of framework. They train the generator and discriminator through an adversarial process.

D. The input of the discriminator is mainly sample data provided by the generator.

C. GANS are trained using the BP algorithm.

Congratulations! Correct Answer: AD

PART 3

39. Which of the following functions are supported by the HiAl Engine platform?

C. Video summarization

D. Automatic generation of video thumbnails

B. Keyword extraction

A. Form recognition

Congratulations! Correct Answer: ABCD

40. Which of the following statements are true about grid search based on hyperparameter tuning?

B. Grid search is expensive and time-consuming.

C. Grid search works well when there are relatively few hyperparameters.

D. Grid search suits neural networks well.

A. Grid search exhaustively searches for all possible hyperparameter combinations to form a hyperparameter value grid.

Congratulations! Correct Answer: ABC

41. Which of the following statements are true about the MindSpore components?

C. nn: neural network cells in MindSpore that define loss functions and optimizers.

B. dataset: dataset processing module that loads and preprocesses data.

A. communication: processes data flow communication between the CPU and memory.

D. train: relates to training model and model quantization module.

Congratulations! Correct Answer: BC

42. Which of the following are supported by data management on ModelArts?

OD. Automatic data migration

C. Intelligent data filtering

B. Team labeling

A. A wide range of data formats

Congratulations! Correct Answer: ABC

43, Backed by HUAWEI CLOUD's accumulated knowledge and expertise in Al, big data, and other cutting-edge technologies, GeoGenius offers a one-stop Al development cloud platform for remote sensing data processing, mining, and management. Which of the following are the main functions of GeoGenius?

B. Afforestation

A. Weather forecasting

C. Emergency response and disaster prevention

D. Smart transport infrastructure

Sorry! Correct Answer: AC

44. Which of the following format conversion operations are not performed by the storage control unit of the Da Vinci Architecture?

B. Video stream compression

D. Img2Col

A. Zero padding and transposition

C. Color space conversion

Congratulations! Correct Answer: BC

45. Which of the following are common ensemble learning algorithms in machine learning?

C. Polynomial regression

B. GBDT

A. Random forest

D. Adaboost

Congratulations! Correct Answer: ABD

46. Which of the following products can be equipped with Ascend 310 Processors?

B. Atlas 500 Al edge station

D. Atlas 300 Al accelerator card

A. Atlas 200 DK Al developer kit

C. Atlas 200 Al accelerator module

Congratulations! Correct Answer: ABCD

47. Which of the following use the Da Vinci Architecture?

A. Ascend 310

B. RTX3080

C. Kunpeng 920

D. Ascend 910

Congratulations! Correct Answer: AD

48. Which of the following are model hyperparameters?

A. Learning rate, iteration count, and batch size in a training neural network

B. K in the k-nearest neighbors (KNN) algorithm

C. Number of trees in a random forest

D. Weight coefficient w in linear regression

Congratulations! Correct Answer: ABC

49. Which of the following statements are true about hidden layers?

A. The more hidden layers, the stronger the identification capability.

D. Insufficient hidden layers may cause overfitting.

B. The more hidden layers, the better.

C. Excessive hidden layers may cause overfitting.

Sorry! Correct Answer: AC

50. Which of the following statements are false about the functions of the pooling layers?

C. Pooling reduces the size of the input data at the next layer and the number of parameters, but increases the computation amount.

B. Max pooling achieves invariance to some extent.

D. Pooling helps prevent overfitting.

A. They reduce the receptive field.

Congratulations! Correct Answer: AC

51. Which of the following statements are false about TensorFlow?

C. TensorFlow 2.0 uses the dynamic graph mechanism by default, with a higher running efficiency than the static graph mechanism.

D. TensorFlow 1.0 uses the graph-session mechanism.

A. TensorFlow can only be used in a deep learning algorithm.

B. TensorFlow supports only GPUs and CPUs.

Sorry! Correct Answer: ABC

52. Which of the following statements are true about decision trees?

A. A decision tree is a binary tree structure.

D. Building a decision tree means selecting and measuring feature attributes and determining their topology.

C. Except the root node, all nodes in a decision tree are called leaf nodes.

B. Each non-leaf node in a decision tree represents a test on a feature attribute; each branch represents the output of the feature attribute within a value range; and each leaf node holds a class label.

Congratulations! Correct Answer: BD

53. Which of the following statements are false about the Gated Recurrent Unit (GRU)?

B. GRU is a variant of a Convolutional Neural Network (CNN).

D. GRU assembles the output of the previous time step with the input of the current time step.

A. Unlike long short-term memory (LSTM), GRU merges the cell state and hidden state.

C. GRU combines the forget and update gates into a single input gate.

Sorry! Correct Answer: BC

54. Which of the following are steps of the Back Propagation Trough Time (BPTT) algorithm?

A. Calculating the output of each neuron (forward)

D. Calculating the layers of neurons

B. Calculating the error of each neuron (backward)

C. Calculating the gradient of each weight

Congratulations! Correct Answer: ABC

PART 4

39. Which of the following products can be equipped with Ascend 310 Processors?

A. Atlas 200 DK Al developer kit

C. Atlas 200 Al accelerator module

D. Atlas 300 Al accelerator card

B. Atlas 500 Al edge station

Congratulations! Correct Answer: ABCD

40, Which of the following statements are true about regression analysis?

A. Regression analysis is a statistical analysis method used to determine the quantitative relationship between two or more variables.

C. Regression analysis is a type of supervised learning.

OD. Linear regression with an absolute loss (L2 regularization) is called Lasso regression.

B. Regression analysis is a type of unsupervised learning.

Congratulations! Correct Answer: AC

41. Which of the following are model hyperparameters?

B. K in the k-nearest neighbors (KNN) algorithm

C. Number of trees in a random forest

A. Learning rate, iteration count, and batch size in a training neural network

OD. Weight coefficient w in linear regression

Congratulations! Correct Answer: ABC

42. Backed by HUAWEI CLOUD's accumulated knowledge and expertise in AI, big data, and other cutting-edge technologies, GeoGenius offers a one-stop AI development cloud platform for remote sensing data processing, mining, and management. Which of the following are the main functions of GeoGenius?

C. Emergency response and disaster prevention

D. Smart transport infrastructure

B. Afforestation

A. Weather forecasting

Congratulations! Correct Answer: AC

43. Which of the following statements are true about common activation functions in deep learning?

A. The sigmoid function is monotonic, continuous, and easy to derive. Its output is bounded, making the network converge better.

D. During training of a deep neural network, the sigmoid, tanh, and softsign functions cannot prevent the vanishing gradient problem.

C. The tanh function is symmetric with respect to the origin, and the mean of its output is closer to 0.

B. The sigmoid function can easily cause the gradient to explode.

Sorry! Correct Answer: CD

44. Which of the following are major advantages of Keras?

OB. Low-level coding style

C. Modular and combinable

A. Easy to use

D. Easy to scale

Congratulations! Correct Answer: ACD

45. Which of the following are used for backpropagation?

B. Activation functions

D. Neuron weight

C. Number of neurons at a single layer

OA. Number of network layers

Congratulations! Correct Answer: BD

46. Which of the following are the features of the Speech Interaction Service (SIS) provided by HUAWEI CLOUD

EI?

A. Multiple (more than three) languages

C. Low requirements

B. High usability

D. Low labor cost

Congratulations! Correct Answer: BCD

47. Which of the following use the Da Vinci Architecture?

A. Ascend 310

OC. Kunpeng 920

OB. RTX3080

D. Ascend 910

Congratulations! Correct Answer: AD

48. Which of the following format conversion operations are not performed by the storage control unit of the Da Vinci Architecture?

OD. Img2Col

C. Color space conversion

A. Zero padding and transposition

B. Video stream compression

Congratulations! Correct Answer: BC

49. Which of the following statements are true about decision trees?

B. Each non-leaf node in a decision tree represents a test on a feature attribute; each branch represents the output of the feature attribute within a value range; and each leaf node holds a class label.

D. Building a decision tree means selecting and measuring feature attributes and determining their topology.

C. Except the root node, all nodes in a decision tree are called leaf nodes.

A. A decision tree is a binary tree structure.

Congratulations! Correct Answer: BD

50. Which of the following statements are true about the three main schools of AI?

A. Behaviorism states that intelligence requires no knowledge, representation, or reasoning.

C. Behaviorism states that intelligent behavior can only be expressed through constant interaction between humans and the surrounding environment in the real world.

D. Symbolism states that Al focuses on behavior control, adaptive computation, and evolutionary computation.

B. Connectionism states that Al can evolve like human intelligence.

Sorry! Correct Answer: AC

51. Which of the following statements are true about activation functions?

B. If we do not use an activation function, the output signals will be just a simple linear function.

D. Activation functions play an important role in learning and understanding complex and nonlinear functions of neural network models.

A. The existence of the activation function introduces linearity into the network.

C. There are many activation functions. You need to select one based on the actual situation.

Congratulations! Correct Answer: BCD

52. Which of the following are steps of the Back Propagation Trough Time (BPTT) algorithm?

B. Calculating the error of each neuron (backward)

C. Calculating the gradient of each weight

D. Calculating the layers of neurons

A. Calculating the output of each neuron (forward)

Congratulations! Correct Answer: ABC

53. Which of the following operations are involved when TensorFlow is used for model training?

C. Obtain and cleanse data.

B. Train and save models.

A. Define input and output nodes.

D. Define learning parameters and operations.

Congratulations! Correct Answer: ABD

54. Which of the following are supported by data management on ModelArts?

A. A wide range of data formats

OD. Automatic data migration

B. Team labeling

C. Intelligent data filtering

Congratulations! Correct Answer: ABC

PART 5

39. Which of the following are supported by data management on ModelArts?

B. Team labeling

D. Automatic data migration

C. Intelligent data filtering

A. A wide range of data formats

答错了!正确答案: ABC

40. Which of the following statements are false about the universal engine of the Ascend 310 software stack?

B. Provides compute resources and executes specific computing tasks.

C. Provides compute resources for Ascend AI.

D. Outputs tasks at the operator layer for hardware.

A. Provides general capabilities needed for neural network training.

恭喜你,答对了! 正确答案: ABCD

41. Which of the following are common tensor operations in MindSpore?

B. astype()

D. abs()

A. dtype()

OC. switch()

恭喜你,答对了!正确答案: ABD

42. Which of the following statements are true about MindSpore and Huawei all-scenario solutions?

B. MindSpore Serving supports batching.

C. Distributed training of ultra-large models and ultra-large datasets requires only data parallelism.

D. MindIR is the key to "Train Once, Deploy Everywhere."

A. The full-scenario deployment solution includes model generation and efficient execution.

恭喜你,答对了!正确答案:ABD

43. Which of the following format conversion operations are not performed by the storage control unit of the Da Vinci Architecture?

A. Zero padding and transposition

D. Img2Col

C. Color space conversion

B. Video stream compression

恭喜你,答对了!正确答案:BC

44. What are the theories of cognitive intelligence?

OC. Behavioral theory

B. The theory of evolution

A. Threshold theory

D. Logic theory

恭喜你,答对了!正确答案:ABD

45. Which of the following development modes are supported by the ModelArts training platform?

C. On-premises development (PyCharm+PyCharm Toolkit)

D. On-premises development (IDE+SDK)

A. In-cloud development (CodeLab)

B. In-cloud development (Notebook+SDK)

恭喜你,答对了! 正确答案: ABCD

46. Which of the following statements are false about the combination of the model bias and variance?

C. A high bias and low variance may result in underfitting.

D. A model with a high variance is of no application value.

A. A good model requires a high variance and high bias.

B. A high bias and high variance may result in overfitting.

恭喜你,答对了!正确答案:ABD

47. Which of the following operations are involved when TensorFlow is used for model training?

C. Obtain and cleanse data.

D. Define learning parameters and operations.

B. Train and save models.

A. Define input and output nodes.

恭喜你,答对了!正确答案:ABD

48. Which of the following statements are false about common optimizers?

B. When the momentum optimizer is used, parameters are updated by using the same learning rate, but momentum coefficients keep changing with each iteration.

C. Adam sets different learning rates for different parameters.

A. One of the advantages of Adagrad optimizers is that the parameter update operation does not end too early.

D. The RMSprop optimizer introduces an attenuation coefficient to ensure a certain proportion of gradient attenuation in each round.

恭喜你,答对了!正确答案:AB

49. Which of the following statements are true about support vector machines (SVMs)?

C. The learning algorithm of SVMs is the optimal algorithm for concave quadratic programming.

D. Based on the structural risk minimization (SRM), SVMs build an optimal hyperplane in the feature space so that the learner can be globally optimized.

A. SVMs are binary classification models. Their basic model is the linear classifier with the largest interval defined in the feature space.

B. SVMs also have a kernel trick, which enables them to perform as a nonlinear classifier.

恭喜你,答对了!正确答案:ABCD

50. Which of the following are common application scenarios of HUAWEI CLOUD OCR?

C. VAT invoice recognition

B. Driving license recognition

D. Aircraft itinerary recognition

A. Photo recognition for cats and dogs

恭喜你,答对了!正确答案:BCD

51. Which of the following statements are true about the MindSpore components?

C. nn: neural network cells in MindSpore that define loss functions and optimizers.

A. communication: processes data flow communication between the CPU and memory.

D. train: relates to training model and model quantization module.

B. dataset: dataset processing module that loads and preprocesses data.

恭喜你,答对了!正确答案:BC

52. Which of the following statements are true about gradient descent?

C. Mini-batch gradient descent (MBGD) is a balance between BGD and SGD, and is the optimal choice for all datasets.

B. Batch gradient descent (BGD) is the most unstable method and consumes too many compute resources.

A. Stochastic gradient descent (SGD) randomly chooses samples for each training job, causing unstability. As a result, the loss function fluctuates or even encounters reverse displacements during the process of dropping to the minimum value.

D. BGD uses all samples for each training job.

恭喜你,答对了!正确答案: ACD

53. Which of the following functions are supported by the HIAI Engine platform?

D. Automatic generation of video thumbnails

C. Video summarization

B. Keyword extraction

A. Form recognition

恭喜你,答对了!正确答案:ABCD

54. Which of the following statements are true about ensemble learning?

C. A batch of features are randomly selected for the subtree training in a random forest.

D. Xgboost supports parallel training like the random forest.

A. Ensemble learning uses the decision tree as the bottom-layer algorithm.

B. The voting policy of ensemble learning can be the averaging method.

恭喜你,答对了!正确答案:BC

BLANK FILLING

The derivative of the ReLU activation function in the negative half interval is fixed at (). (Fill in the blank with a number.)

0

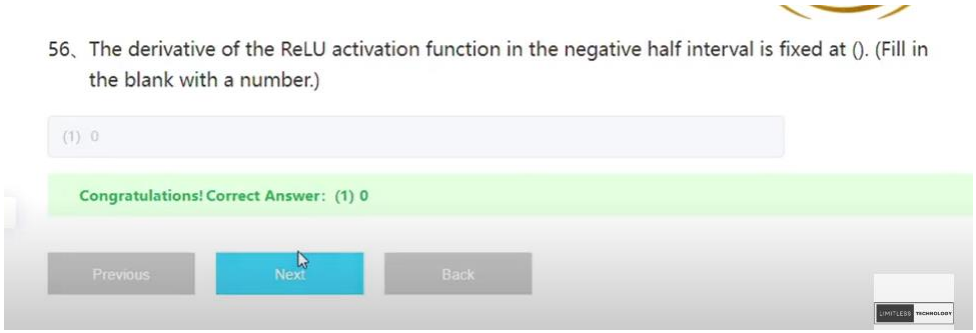

Tensor [[[2,3]]] is a/an ()-dimensional tensor. (Fill in the blank with a number.)

3

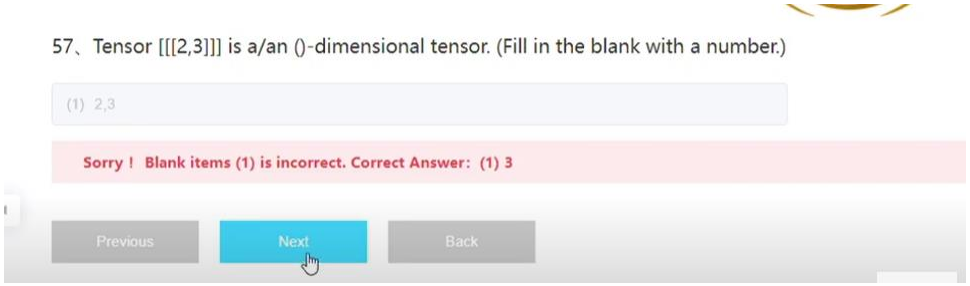

As the neural network increasingly deepens, two issues may be encountered during network training: one is (), and the other is gradient vanishing/exploding.

Network degradation;network degradation:Network Degradation

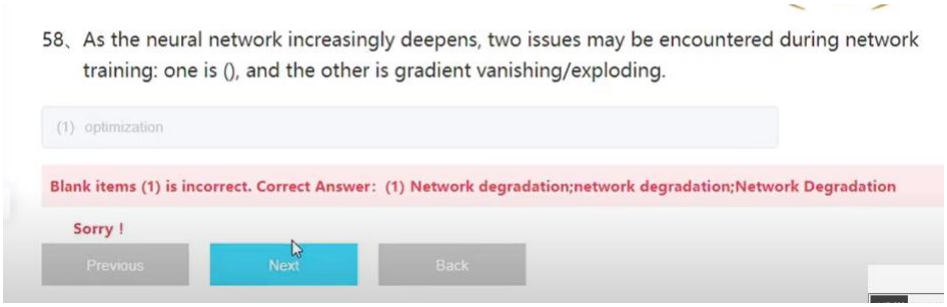

The average value of tensors generated using tf.random.normal([4], 2, 1, tf.float32) is (). (Fill in the blank with a number.)

2

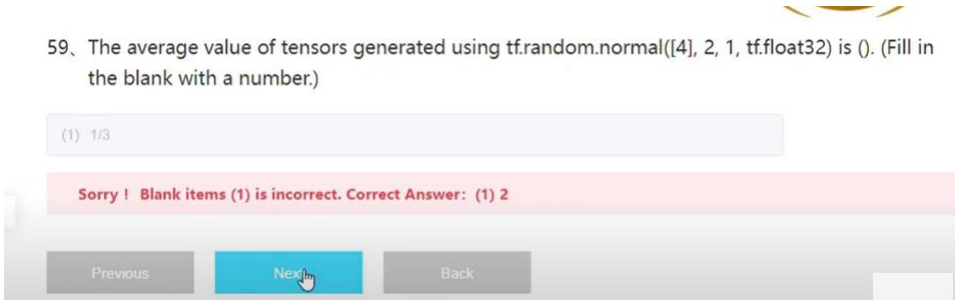

Generally, a model with a () volume has a higher precision, and a model with a () volume has a higher efficiency. (Fill in using "large" or "small")

(1) large (2) small

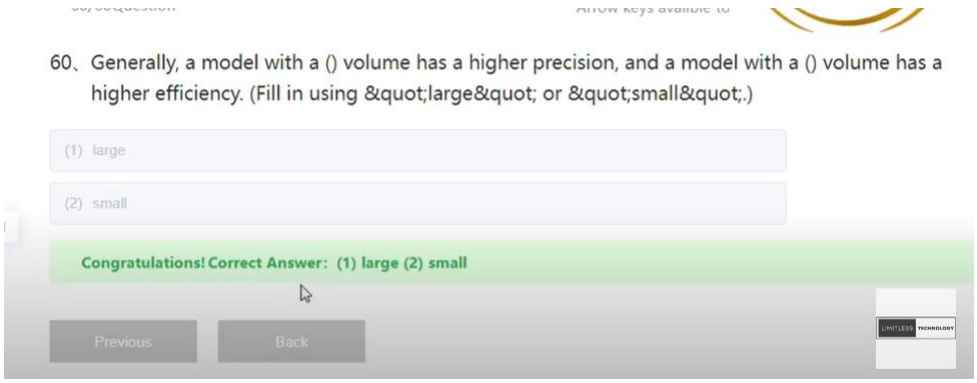

Parts : 2

Generally, a model with a () volume has a higher precision, and a model with a () volume has a higher efficiency. (Fill in using "large" or "small".)

(1) large (2) small

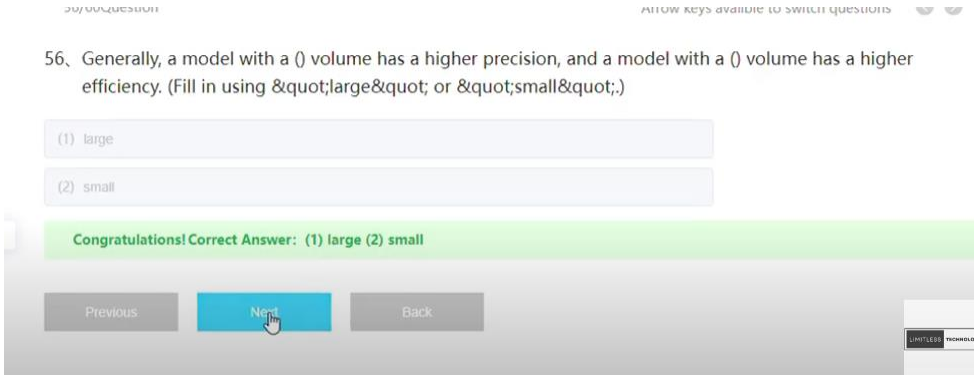

The Naive Bayes algorithm needs to obtain the () probability of the dataset.

PRIOR

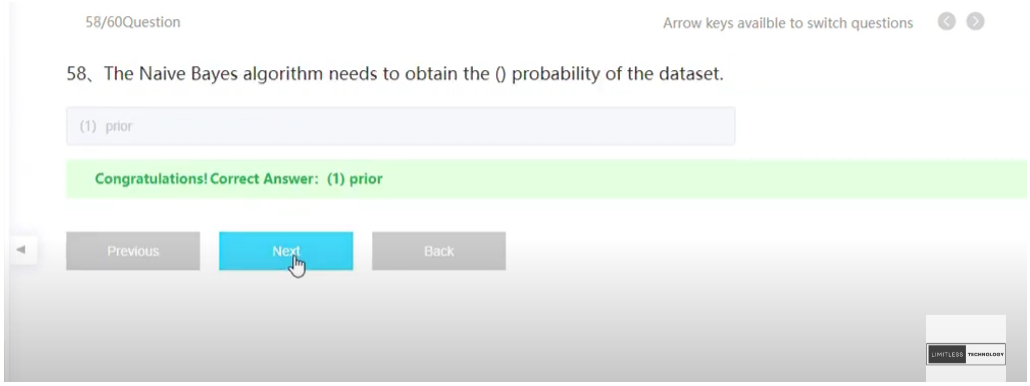

The 0 subsystem of MindSpore provides Mind IR-oriented, graph-level, real-time compilation.

(1) MindCompiler

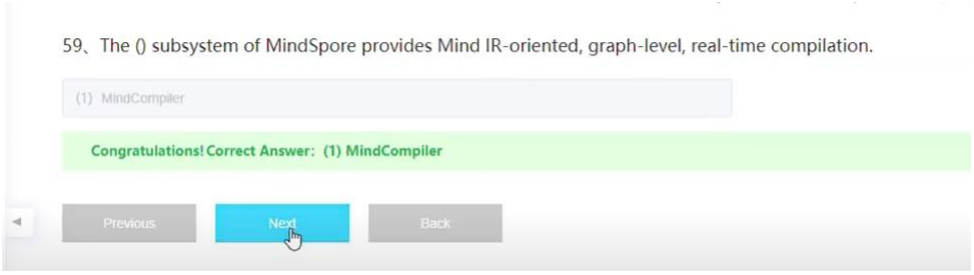

Parts 3 ;

As the neural network increasingly deepens, two issues may be encountered during network training: one is 0, and the other is gradient vanishing/exploding.

Network degradation network degradation; Network Degradation

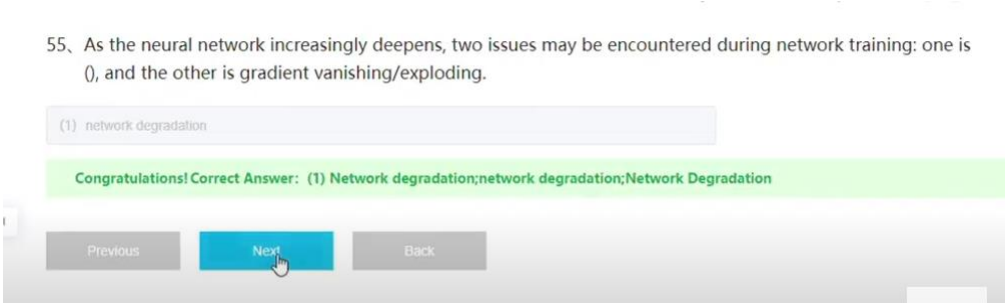

Linear regression using a loss function with the LO-norm regular term is also called ridge regression.
2

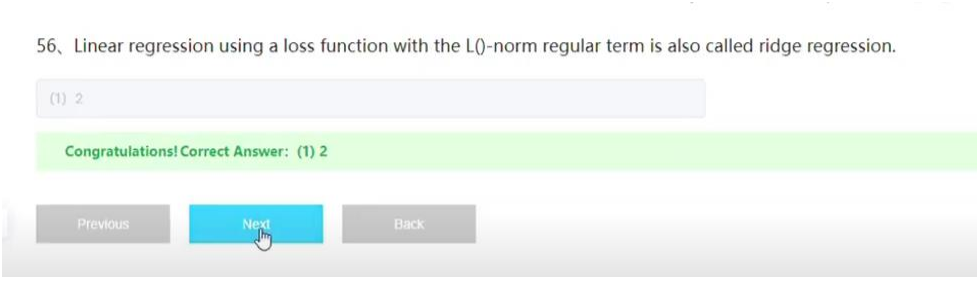

.Assume that there are 10,000 data pieces in a cancer data sample, of which 100 pieces are from cancer patients, and the other 9900 pieces are normal. If a classification model predicts that 9000 out of the 9900 pieces are normal, and 90 out of the 100 pieces are from cancer patients, the accuracy rate of the model is 0%.

90.9

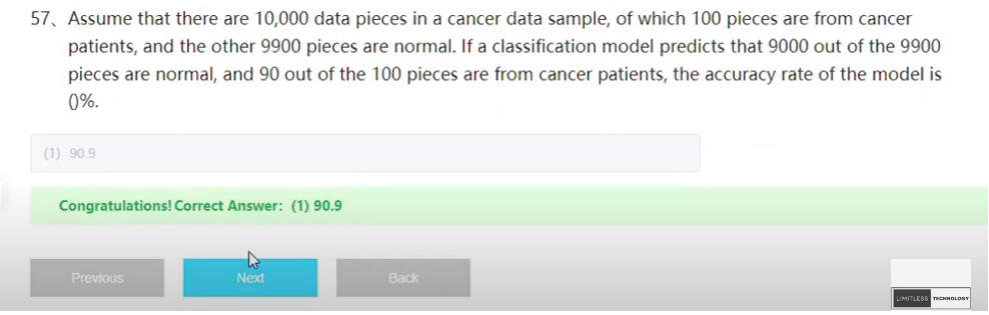

The Atlas 200 DK is designed to run on the Ascend () AI Processor. (Full in using Arabic numerals.)

(1) 310

Congratulations! Correct Answer: (1) 310

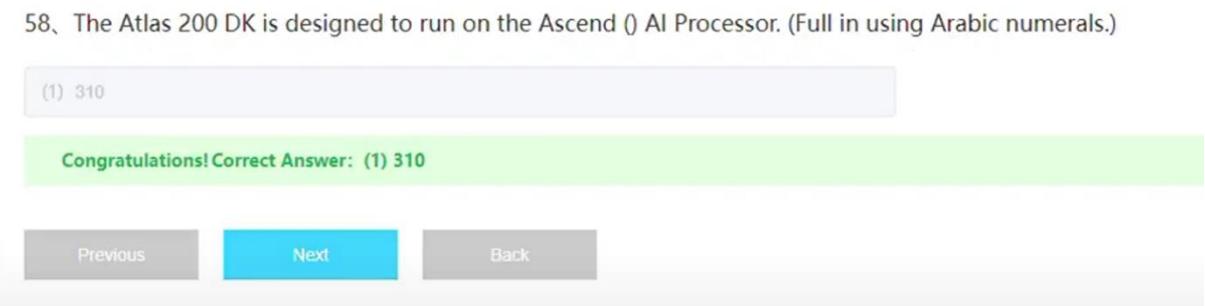

Generally, a model with a () volume has a higher precision, and a model with a () volume has a higher efficiency. (Fill in using "large" or "small")

(1) large (2) small

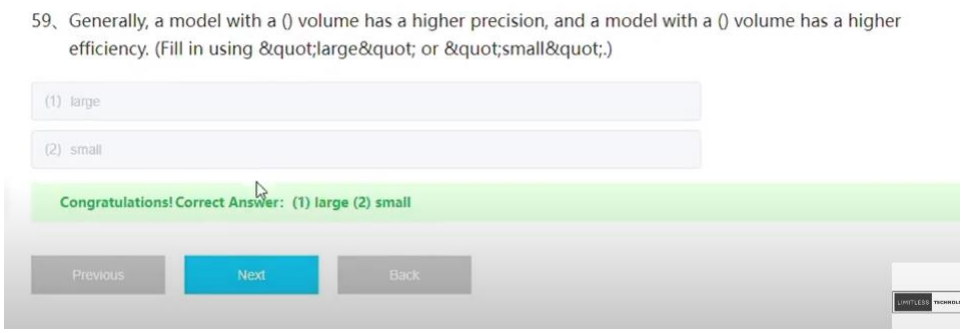

The () algorithm (abbreviation) is used when RNN performs backpropagation (BP).

(1) BPTT

Congratulations! Correct Answer: (1) BPTT;bptt

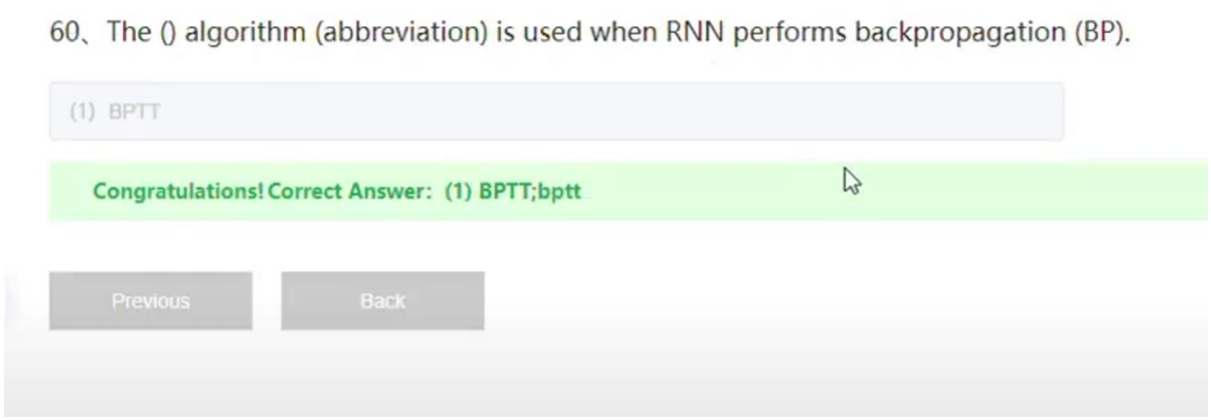

Prt 4

The positive () direction is the maximum directional derivative of a function.

Congratulations! Correct Answer: (1) gradient

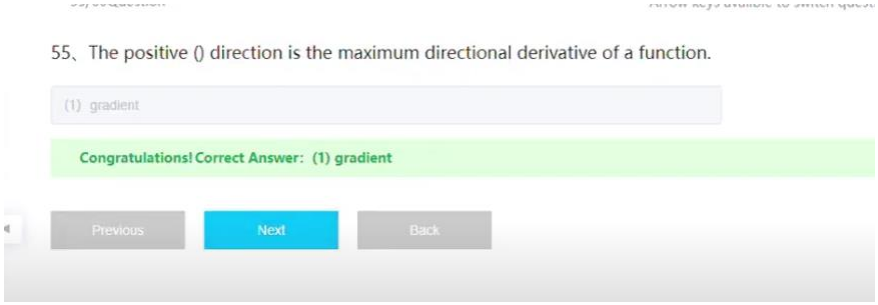

Tensor [[[2,3]]] is a/an ()-dimensional tensor. (Fill in the blank with a number.)

3

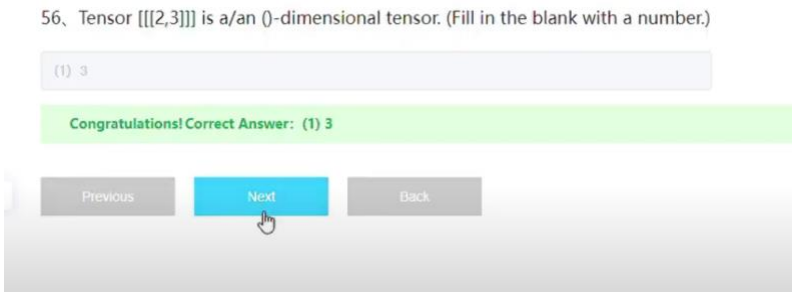

In TensorFlow 2.0, if tf.keras.layers.RNN is used to process timing information and you want to obtain the output status at each moment, set () to True.

Congratulations! Correct Answer: (1) return_sequence

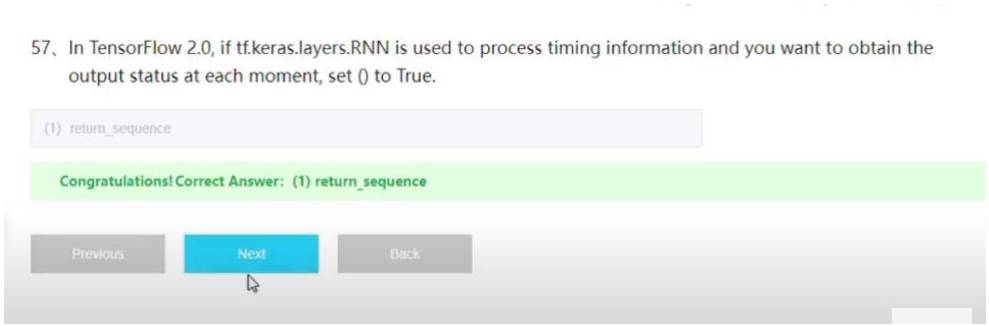

The average value of tensors generated using tf.random.normal([4], 2, 1, tf.float32) is (). (Fill in the blank with a number.)

(1):2

Congratulations! Correct Answer: (1) 2

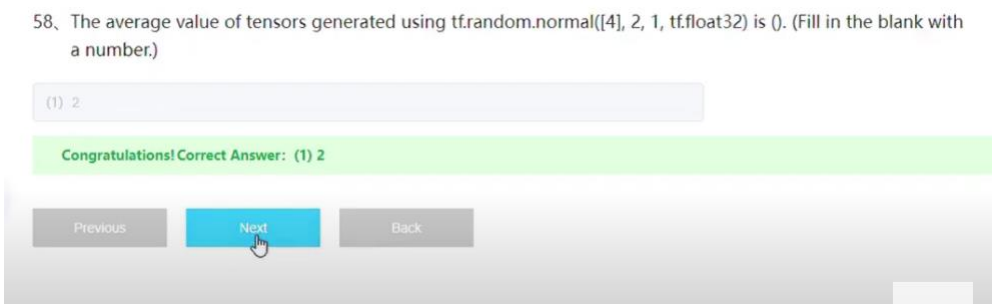

mindspore.ops.GradOperation is the first-order derivative method in MindSpore to compute gradients. It contains multiple parameters. When get all is set to 0), the derivatives of all parameters are computed.

(1) bue

Sorry! Blank items (1) is incorrect. Correct Answer: (1) True

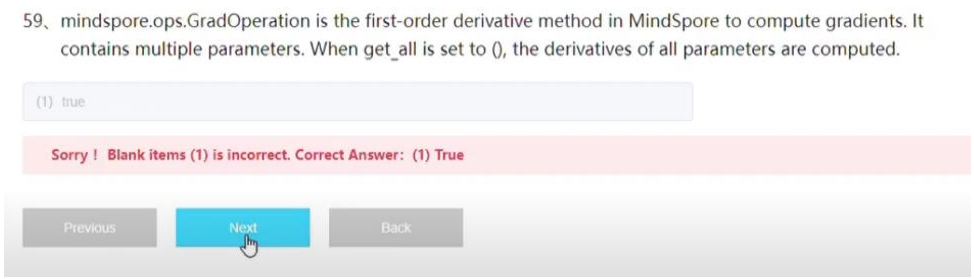

Sensors and loT technologies provide technical support for the () of the four elements of Al.

(1) dutx

Congratulations! Correct Answer: (1) data;Data;DATA

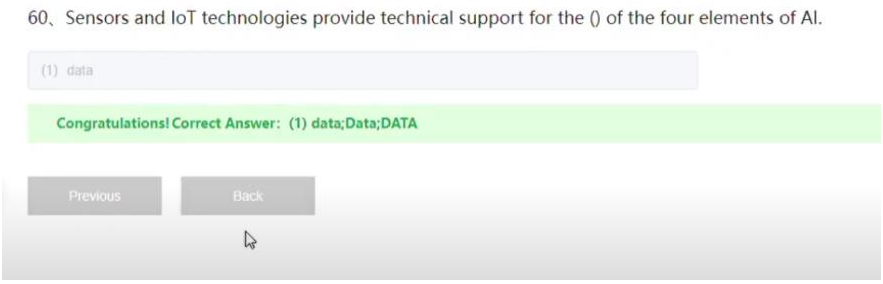

PARTS : 5

The Speech Interaction Service (SIS) on HUAWEI CLOUD provides text recognition through open (). (Fill in the blank with the abbreviation.)

(1) APIs

恭喜你,答对了!正确答案: (1) APIs

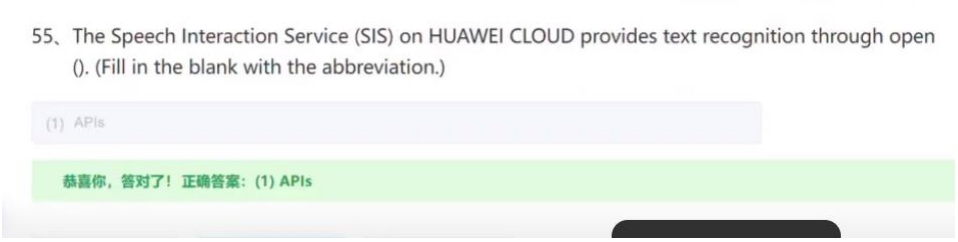

Generally, a model with a () volume has a higher precision, and a model with a () volume has a higher efficiency. (Fill in using "large" or "small".)

(1) large

(2) amall

恭喜你,答对了!正确答案: (1) large (2) small

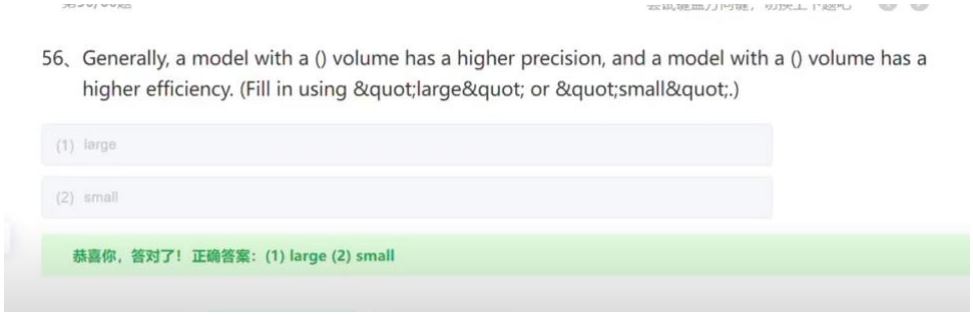

In TensorFlow 2.0, if tf.keras.layers.RNN is used to process timing information and you want to obtain the output status at each moment, set () to True.

(1) tum querie

恭喜你,答对了!正确答案: (1) return sequence

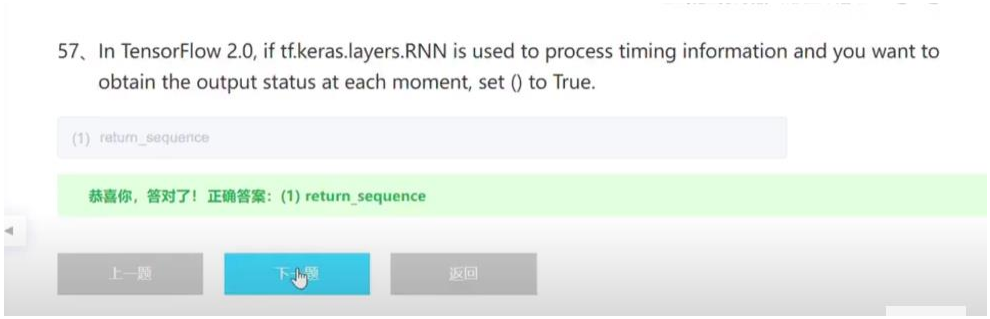

At a convolution layer, there are 256 5 x 5 convolution kernels, the size of an input feature map is 32 x 32 x 5, the convolution step is 2, and zero padding is employed (padding size=1). So, the size of the output feature map is ()*()*(). (Fill in each blank with a number.)

(1) 15

(2) 15

(3) 256

答错了! 填空(1) (2) (3) 错误:正确答案:

(1) 15

(2) 15

(3) 256